

Deep Learning Lab 5 - Report

([Github Link](#))

Training Loss and Convergence Analysis

We implemented three models for next-word prediction using a dataset of 100 poems:

1. A Recurrent Neural Network (RNN) implemented from scratch using NumPy.
2. A PyTorch RNN using One-Hot Encoding.
3. A PyTorch RNN using Trainable Word Embeddings.

All models were trained to predict the next word in a sequence of five input words. The training loss was monitored across epochs to evaluate convergence behavior and stability.

The NumPy-based RNN exhibited the slowest convergence. Due to manual implementation of forward and backward propagation, the loss decreased gradually and remained comparatively higher. The absence of optimized tensor operations and automatic differentiation limited its efficiency.

The One-Hot Encoding model demonstrated improved convergence behavior. The loss decreased steadily across epochs without major fluctuations, indicating stable optimization. However, the input dimensionality was equal to the vocabulary size, resulting in high-dimensional sparse vectors that increased computational overhead.

The Embedding-based RNN achieved the fastest convergence and the lowest final loss. Since words were represented as dense trainable vectors, the model captured semantic relationships more effectively. The loss curve showed smooth and stable decline, confirming efficient gradient updates and better representation learning.

Training Time and Final Loss Comparison

The table below summarizes the performance of all three models:

Model	Training_Time_Seconds	Final_Loss
RNN From Scratch (NumPy)	11621.34060907360	22.814732670355100
One-Hot Encoding (PyTorch)	257.89110493660000	0.0331530771172118
Trainable Embedding (PyTorch)	130.76644897460900	0.038799241625699

The scratch implementation required the longest training time due to manual gradient computation. The One-Hot model trained faster because of PyTorch's optimized backend. The Embedding model showed comparable or slightly faster training time while achieving lower loss, demonstrating better efficiency.

Evaluation of Generated Text Quality

The quality of generated text was evaluated qualitatively by observing coherence, grammatical structure, and contextual consistency.

The NumPy-based model produced text that often contained repetitive word patterns and limited contextual flow. This behavior is expected due to limited optimization capacity and shorter training duration.

The One-Hot model generated grammatically structured sequences but frequently repeated common words. Since one-hot vectors do not encode semantic similarity between words, the model treated each word independently, limiting contextual understanding.

The Embedding-based model produced noticeably more coherent and contextually meaningful sentences. Because embeddings capture semantic proximity between words, the generated text demonstrated smoother transitions and improved sentence structure.

Overall, the embedding approach yielded the most natural and meaningful text generation results.

Advantages and Disadvantages

1. RNN From Scratch

Advantages:

- Provides deep understanding of RNN internal mechanics.
- Full control over gradient computation and parameter updates.

Disadvantages:

- Computationally inefficient.
- Slower convergence.
- Higher implementation complexity.
- More prone to numerical instability.

2. One-Hot Encoding

Advantages:

- Simpler implementation using PyTorch.
- Stable convergence.
- Easy to interpret input representation.

Disadvantages:

- High-dimensional sparse vectors.
- No semantic similarity between words.
- Increased memory consumption.

3. Trainable Embeddings

Advantages:

- Dense and compact word representations.
- Captures semantic relationships.
- Faster convergence.
- Lower final loss.
- Better text quality.

Disadvantages:

- Requires tuning embedding dimension.
- Slightly more complex architecture than one-hot.

Overall Conclusion

This experiment demonstrates that word representation plays a critical role in sequence modeling tasks. While the scratch implementation is valuable for conceptual understanding, it is not computationally efficient for practical applications. The One-Hot Encoding approach improves optimization stability but lacks semantic awareness. The Trainable Embedding approach achieves superior performance in terms of training efficiency, lower loss, and improved text coherence.

Thus, embedding-based representations are more suitable for real-world text generation tasks using recurrent neural networks.